

Pr. 1.

For $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$:

$$\mathbf{B} = \begin{bmatrix} \mathbf{0} & \mathbf{A} \\ \mathbf{A}^T & \mathbf{0} \end{bmatrix} = \begin{bmatrix} \mathbf{0} & \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T \\ \mathbf{V}\mathbf{\Sigma}\mathbf{U}^T & \mathbf{0} \end{bmatrix} \quad (\mathbf{\Sigma}^T = \mathbf{\Sigma} \text{ since } \mathbf{A} \text{ is square}).$$

Hence

$$\mathbf{B} \begin{bmatrix} u_i \\ v_i \end{bmatrix} = \begin{bmatrix} \mathbf{0} & \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T \\ \mathbf{V}\mathbf{\Sigma}\mathbf{U}^T & \mathbf{0} \end{bmatrix} \begin{bmatrix} u_i \\ v_i \end{bmatrix} = \begin{bmatrix} \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T v_i \\ \mathbf{V}\mathbf{\Sigma}\mathbf{U}^T u_i \end{bmatrix} = \begin{bmatrix} \sigma_i u_i \\ \sigma_i v_i \end{bmatrix} = \sigma_i \begin{bmatrix} u_i \\ v_i \end{bmatrix}.$$

Thus $\mathbf{x} \triangleq \begin{bmatrix} u_i \\ v_i \end{bmatrix}$ is an eigenvector of \mathbf{B} corresponding to the eigenvalue σ_i . Similarly:

$$\mathbf{B} \begin{bmatrix} u_i \\ -v_i \end{bmatrix} = \begin{bmatrix} \mathbf{0} & \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T \\ \mathbf{V}\mathbf{\Sigma}\mathbf{U}^T & \mathbf{0} \end{bmatrix} \begin{bmatrix} u_i \\ -v_i \end{bmatrix} = \begin{bmatrix} -\mathbf{U}\mathbf{\Sigma}\mathbf{V}^T v_i \\ \mathbf{V}\mathbf{\Sigma}\mathbf{U}^T u_i \end{bmatrix} = \begin{bmatrix} -\sigma_i u_i \\ \sigma_i v_i \end{bmatrix} = -\sigma_i \begin{bmatrix} u_i \\ -v_i \end{bmatrix},$$

so that $\mathbf{z} \triangleq \begin{bmatrix} u_i \\ -v_i \end{bmatrix}$ is an eigenvector of \mathbf{B} corresponding to the eigenvalue $-\sigma_i$. Note that $\mathbf{z}'\mathbf{x} = [u_i^T \quad -v_i^T] \begin{bmatrix} u_i \\ v_i \end{bmatrix} = 0$ so \mathbf{x} and \mathbf{z} are orthogonal. Thus the eigenvalues of \mathbf{B} are exactly $\{\pm\sigma_i\}$.

We must normalize the eigenvectors to have unit norm. The unit-norm eigenvector of \mathbf{B} corresponding to the eigenvalue σ_i is $\begin{bmatrix} u_i \\ v_i \end{bmatrix} / \sqrt{2}$ (check that this is unit norm!) and the unit-norm eigenvector corresponding to the eigenvalues $-\sigma_i$ is $\begin{bmatrix} u_i \\ -v_i \end{bmatrix} / \sqrt{2}$, where u_i and v_i are the left and right singular vectors of \mathbf{A} associated with the singular value σ_i .

Pr. 2.

(a) Here $\mathbf{A} = \mathbf{x}\mathbf{y}^T$ is written in outer product form. We can express \mathbf{A} as

$$\mathbf{A} = \mathbf{x} \begin{bmatrix} y_1 & \dots & y_n \end{bmatrix} = \begin{bmatrix} y_1 \mathbf{x} & \dots & y_n \mathbf{x} \end{bmatrix},$$

which has (at most) one linearly independent column because for every $i \neq j$, if $\mathbf{A}_{:,i}$ and $\mathbf{A}_{:,j}$ denote the i th and j th column of \mathbf{A} , then

$$y_i \mathbf{A}_{:,j} - y_j \mathbf{A}_{:,i} = y_i y_j \mathbf{x} - y_j y_i \mathbf{x} = 0.$$

Because $\mathbf{y} \neq \mathbf{0}$ is given, we can choose i and j such that at least one of y_i or y_j is nonzero; thus the columns of \mathbf{A} are linearly dependent. Because $\mathbf{x}, \mathbf{y} \neq \mathbf{0}$ is given, \mathbf{A} has rank one.

(If either of \mathbf{x} or \mathbf{y} were zero, then \mathbf{A} would have rank zero.)

We can see this another way via the SVD. Write $\mathbf{A} = \sigma \mathbf{u}\mathbf{v}^T$ where $\sigma \triangleq \sqrt{\langle \mathbf{x}, \mathbf{x} \rangle \cdot \langle \mathbf{y}, \mathbf{y} \rangle}$, $\mathbf{u} \triangleq \mathbf{x} / \sqrt{\langle \mathbf{x}, \mathbf{x} \rangle}$ and $\mathbf{v} \triangleq \mathbf{y} / \sqrt{\langle \mathbf{y}, \mathbf{y} \rangle}$. Comparing this with the SVD form $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$ we have that the first column of \mathbf{U} and \mathbf{V} are equal to the \mathbf{u} and \mathbf{v} derived. The remaining columns of \mathbf{U} and \mathbf{V} can be any set of orthogonal vectors that is in the ortho-complement of \mathbf{u} and \mathbf{v} , respectively. Further, $\mathbf{\Sigma}$ has one nonzero element, given by σ . Because there is at most one nonzero singular value, the rank of \mathbf{A} is 1, implying \mathbf{A} has one linearly independent column.


(b) Typing the given Julia code yields Figure 1, where there $n = 100$ non-zero singular values. The discrepancy arises because of numerical round-off error due to representing the matrix \mathbf{A} in finite-precision arithmetic. Typing `rank(A)` gives the theoretically correct answer despite the $n - 1$ singular values not being identically equal to zero. This is because Julia's `rank` function first computes the singular values and counts how many are greater than $\max(m, n)\epsilon$ where ϵ is machine precision (equal to 2.204×10^{-16} on my machine).

(c) Typing `@which rank(A)`

```
rank(A::AbstractArray{T<:Any,2}) at linalg\generic.jl:308
```

Clicking on the link to `linalg/generic.jl` reveals the rank function

```
function rank(A::AbstractMatrix)
    m,n = size(A)
    (m == 0 || n == 0) && return 0
    sv = svdvals(A)
    return sum(sv .> maximum(size(A))*eps(sv[1]))
end
```

FIGURE 1. The singular values of $\mathbf{A} = \mathbf{x}\mathbf{y}^T$.

This reveals that the 'numerical rank' of the matrix \mathbf{A} is determined by the formula

$$r = \{\text{Number of } \sigma_i > \epsilon \max(m, n)\},$$

where ϵ is machine precision. Note the 'continuous' looking spectrum of the noise-only singular values in Figure 1. This is the signature of noise-only singular values.

For a 1×2 array in double precision (Float64), the tolerance is $2\epsilon = 4.408 \times 10^{-16}$ on a typical 64-bit computer.

Pr. 3.

Let $\mathbf{A} \in \mathbb{R}^{n \times n}$ be a square matrix. The SVD of \mathbf{A} always exists and is given by $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$. When the rank of \mathbf{A} is n , $\mathbf{\Sigma}$ is an $n \times n$ diagonal matrix with strictly positive entries; otherwise it is an $n \times n$ diagonal matrix with $r < n$ strictly positive entries and is hence positive semi-definite. Since \mathbf{V} is an orthogonal matrix, $\mathbf{V}^T\mathbf{V} = \mathbf{I}$ and hence we can express \mathbf{A} as:

$$\begin{aligned} \mathbf{A} &= \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T \\ &= \mathbf{U}\mathbf{I}\mathbf{\Sigma}\mathbf{V}^T \quad \text{Because } \mathbf{A} \text{ is square, } \mathbf{U}, \mathbf{V} \text{ and } \mathbf{I} \text{ have the same dimensions.} \\ &= \underbrace{\mathbf{U}\mathbf{V}^T}_{\triangleq \mathbf{Q}} \underbrace{\mathbf{\Sigma}\mathbf{V}^T}_{\triangleq \mathbf{S}}. \end{aligned}$$

Recall from HW1 that $\mathbf{Q} = \mathbf{U}\mathbf{V}^T$ is an orthogonal matrix because $\mathbf{Q}^T\mathbf{Q} = \mathbf{V}\mathbf{U}^T\mathbf{U}\mathbf{V}^T = \mathbf{V}\mathbf{I}\mathbf{V}^T = \mathbf{I}$. The matrix $\mathbf{S} = \mathbf{V}\mathbf{\Sigma}\mathbf{V}^T$ is symmetric because $\mathbf{S}^T = \mathbf{V}\mathbf{\Sigma}\mathbf{V}^T = \mathbf{S}$. It is positive definite when its eigenvalues, which are exactly equal to the n strictly positive singular values of \mathbf{A} , are strictly positive, and it is positive semi-definite when the rank of \mathbf{A} is less than n . Note $\mathbf{S} = \mathbf{S}^T$ and all the eigenvalues of \mathbf{S} being non-negative (or positive) is a sufficient and necessary test for positive semi-definiteness (resp. positive definiteness) of a matrix.

Pr. 4.

(a) We have that

$$\begin{aligned} \mathbf{A} = \mathbf{x}\mathbf{y}^T &= \frac{\mathbf{x}}{\|\mathbf{x}\|_2} \|\mathbf{x}\|_2 \|\mathbf{y}\|_2 \frac{\mathbf{y}^T}{\|\mathbf{y}\|_2} \\ (1) \quad &= \begin{bmatrix} \frac{\mathbf{x}}{\|\mathbf{x}\|_2} & \mathbf{U}^\perp \end{bmatrix} \begin{bmatrix} \|\mathbf{x}\|_2 \|\mathbf{y}\|_2 & 0 & \dots & 0 \\ 0 & 0 & \dots & 0 \\ \vdots & & \ddots & \\ 0 & 0 & \dots & 0 \end{bmatrix} \begin{bmatrix} \frac{\mathbf{y}}{\|\mathbf{y}\|_2}^T & \mathbf{V}^\perp \end{bmatrix}. \end{aligned}$$

We know that $\mathbf{x}\mathbf{y}^T$ is a rank-1 matrix, and thus has only one nonzero singular value. Furthermore, $\frac{\mathbf{x}}{\|\mathbf{x}\|_2}$ and $\frac{\mathbf{y}}{\|\mathbf{y}\|_2}$ are unit norm vectors, and $\|\mathbf{x}\|_2\|\mathbf{y}\|_2$ is a positive real number. Therefore, $\sigma_1 = \|\mathbf{x}\|_2\|\mathbf{y}\|_2$ is the only nonzero singular value of \mathbf{A} , and $\mathbf{u}_1 = \frac{\mathbf{x}}{\|\mathbf{x}\|_2}$, $\mathbf{v}_1 = \frac{\mathbf{y}}{\|\mathbf{y}\|_2}$ are the corresponding left and right singular vectors. The full SVD of \mathbf{A} is given by (1), where \mathbf{U}^\perp is any set of orthonormal vectors orthogonal to \mathbf{u}_1 , and \mathbf{V}^\perp is any set of orthonormal vectors orthogonal to \mathbf{v}_1 .

Using the definition of pseudoinverse,

$$\mathbf{A}^\dagger = \frac{\mathbf{y}}{\|\mathbf{y}\|_2} \frac{1}{\|\mathbf{x}\|_2\|\mathbf{y}\|_2} \frac{\mathbf{x}^T}{\|\mathbf{x}\|_2} = \frac{\mathbf{y}\mathbf{x}^T}{\|\mathbf{x}\|_2^2\|\mathbf{y}\|_2^2}.$$

(b) This is a special case of (a) in which $\mathbf{y} = \mathbf{x}$. Thus:

$$\mathbf{A}^\dagger = \frac{\mathbf{x}}{\|\mathbf{x}\|_2} \frac{1}{\|\mathbf{x}\|_2\|\mathbf{x}\|_2} \frac{\mathbf{x}^T}{\|\mathbf{x}\|_2} = \frac{\mathbf{x}\mathbf{x}^T}{\|\mathbf{x}\|_2^4}.$$

Pr. 5.

Approach 1:

If $\mathbf{y} \in \mathcal{N}(\mathbf{A}^T)$, then $\mathbf{A}^T\mathbf{y} = \mathbf{0}$, so $\mathbf{x}^T\mathbf{A}^T\mathbf{y} = 0$ for any vector \mathbf{x} , so $(\mathbf{A}\mathbf{x})^T\mathbf{y} = 0$ and $\mathbf{A}\mathbf{x} \perp \mathbf{y}$. Thus $\mathbf{y} \in \mathcal{R}(\mathbf{A})^\perp$ because $\mathcal{R}(\mathbf{A}) = \{\mathbf{A}\mathbf{x}\}$. This shows that $\mathcal{N}(\mathbf{A}^T) \subseteq \mathcal{R}(\mathbf{A})^\perp$.

Conversely, if $\mathbf{y} \in \mathcal{R}(\mathbf{A})^\perp$ then $\mathbf{y} \perp \mathbf{A}\mathbf{x}$ for any \mathbf{x} so $0 = (\mathbf{A}\mathbf{x})^T\mathbf{y} = \mathbf{x}^T(\mathbf{A}^T\mathbf{y})$ so choosing $\mathbf{x} = \mathbf{A}^T\mathbf{y}$ we have $\|\mathbf{A}^T\mathbf{y}\| = 0$ so $\mathbf{y} \in \mathcal{N}(\mathbf{A}^T)$. Thus $\mathcal{N}(\mathbf{A}^T) = \mathcal{R}(\mathbf{A})^\perp$.

Approach 2:

Let $\mathbf{A} \in \mathbb{R}^{m \times n}$. Then $\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^T$ and $\mathbf{A}^T = \mathbf{V}\Sigma^T\mathbf{U}^T$. Using our discussion on the anatomy of the SVD:

$$\begin{aligned} \mathbf{A} &= \begin{bmatrix} \underbrace{\mathbf{u}_1 \dots \mathbf{u}_r}_{\text{basis for } \mathcal{R}(\mathbf{A})} & \underbrace{\mathbf{u}_{r+1} \dots \mathbf{u}_m}_{\text{basis for } \mathcal{R}(\mathbf{A})^\perp} \end{bmatrix} \Sigma \begin{bmatrix} \underbrace{\mathbf{v}_1 \dots \mathbf{v}_r}_{\text{basis for } \mathcal{N}(\mathbf{A})^\perp} & \underbrace{\mathbf{v}_{r+1} \dots \mathbf{v}_n}_{\text{basis for } \mathcal{N}(\mathbf{A})} \end{bmatrix}^T \\ \mathbf{A}^T &= \begin{bmatrix} \underbrace{\mathbf{v}_1 \dots \mathbf{v}_r}_{\text{basis for } \mathcal{R}(\mathbf{A}^T)} & \underbrace{\mathbf{v}_{r+1} \dots \mathbf{v}_n}_{\text{basis for } \mathcal{R}(\mathbf{A}^T)^\perp} \end{bmatrix} \Sigma^T \begin{bmatrix} \underbrace{\mathbf{u}_1 \dots \mathbf{u}_r}_{\text{basis for } \mathcal{N}(\mathbf{A}^T)^\perp} & \underbrace{\mathbf{u}_{r+1} \dots \mathbf{u}_m}_{\text{basis for } \mathcal{N}(\mathbf{A}^T)} \end{bmatrix}^T \end{aligned}$$

We see that $\mathcal{R}(\mathbf{A})^\perp = \text{span}(\{\mathbf{u}_{r+1}, \dots, \mathbf{u}_m\}) = \mathcal{N}(\mathbf{A}^T)$.

Pr. 6.

(a) The problem of finding the line $y = \alpha x + b$ that best fits the points $(1, 2)$, $(2, 1)$ and $(3, 3)$ is equivalent to the

problem: Find β that minimizes $\|\mathbf{z} - \mathbf{A}\beta\|_2$, with $\mathbf{z} \triangleq \begin{bmatrix} 2 \\ 1 \\ 3 \end{bmatrix}$, $\mathbf{A} \triangleq \begin{bmatrix} 1 & 1 \\ 2 & 1 \\ 3 & 1 \end{bmatrix}$, $\beta = \begin{bmatrix} \alpha \\ b \end{bmatrix}$.

Thus the optimal LLS solution is $\hat{\beta} = \mathbf{A}^\dagger \mathbf{z} = \mathbf{V}\Sigma^+ \mathbf{U}^T \mathbf{z} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{z}$.

Here $(\mathbf{A}^T \mathbf{A})^{-1} = \begin{bmatrix} 14 & 6 \\ 6 & 3 \end{bmatrix}^{-1} = \frac{1}{14 \cdot 3 - 6 \cdot 6} \begin{bmatrix} 3 & -6 \\ -6 & 14 \end{bmatrix} = \frac{1}{6} \begin{bmatrix} 3 & -6 \\ -6 & 14 \end{bmatrix}$, and $\mathbf{A}^T \mathbf{z} = \begin{bmatrix} 13 \\ 6 \end{bmatrix}$,

so $\hat{\beta} = \frac{1}{6} \begin{bmatrix} 3 & -6 \\ -6 & 14 \end{bmatrix} \begin{bmatrix} 13 \\ 6 \end{bmatrix} = \frac{1}{6} \begin{bmatrix} 3 \\ 6 \end{bmatrix} = \begin{bmatrix} 0.5 \\ 1 \end{bmatrix}$.

We repeat this calculation using the SVD. Let $\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^T$ denote a SVD of \mathbf{A} . Numerically computing an SVD yields:

$$\mathbf{U} = \begin{bmatrix} -0.3231 & 0.8538 & 0.4082 \\ -0.5475 & 0.1832 & -0.8165 \\ -0.7719 & -0.4873 & 0.4082 \end{bmatrix}, \quad \Sigma = \begin{bmatrix} 4.0791 & 0 \\ 0 & 0.6005 \\ 0 & 0 \end{bmatrix}, \quad \mathbf{V} = \begin{bmatrix} -0.9153 & -0.4027 \\ -0.4027 & 0.9153 \end{bmatrix},$$

so $\hat{\beta} = \mathbf{V}\Sigma^+ \mathbf{U}^T = \mathbf{V} \begin{bmatrix} 0.2451 & 0 & 0 \\ 0 & 1.6653 & 0 \end{bmatrix} \mathbf{U}^T \mathbf{z} = \begin{bmatrix} 0.5 \\ 1 \end{bmatrix}$.

Unsurprisingly we get the same answer since $\mathbf{A}^+ = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}$ when \mathbf{A} has full column rank.

(b) Here we are asked to find the line $y = \alpha x + b$ that best fits the points $(2, 1)$, $(1, 2)$ and $(3, 3)$. We could repeat the above computation, or we can recognize that all that has changed is that the first two points are swapped between (a) and (b). Thus the solution β of the two problems will be identical because the least-squares criterion is the same for any ordering of the (x_i, y_i) points. A more formal way of seeing is to first note that the new \mathbf{A} matrix for (b) is exactly same as the \mathbf{A} matrix in (a) except for the first two rows swapping places. Denote the

“ \mathbf{A} ” matrix in part (a) by \mathbf{A}_a and the “ \mathbf{A} ” matrix in part (b) by \mathbf{A}_b . Then $\mathbf{A}_b = \underbrace{\begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}}_{\triangleq \mathbf{P}} \mathbf{A}_a$.

Note that \mathbf{P} is a orthogonal matrix because $\mathbf{P}\mathbf{P}^T = \mathbf{I}$. Such a \mathbf{P} is called a permutation matrix. We also have that the new $\mathbf{z}_b = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} = \mathbf{P}\mathbf{z}_a$. Thus the new problem is Find β_b that minimizes $\|\mathbf{z}_b - \mathbf{A}_b\beta\|_2 = \|\mathbf{P}\mathbf{z}_a - \mathbf{P}\mathbf{A}_a\beta\|_2 = \|\mathbf{z}_a - \mathbf{A}_a\beta\|_2$, because \mathbf{P} is an orthogonal matrix. Thus we get the same solution β .

Pr. 7.

Define for integer values d :

$$\mathbf{y} \triangleq \begin{bmatrix} f(t_1) \\ \vdots \\ f(t_{16}) \end{bmatrix} \in \mathbb{R}^{16}, \quad \mathbf{A} \triangleq \begin{bmatrix} 1 & t_1 & t_1^2 & \dots & t_1^d \\ \vdots & \vdots & \vdots & \dots & \vdots \\ 1 & t_{16} & t_{16}^2 & \dots & t_{16}^d \end{bmatrix} \in \mathbb{R}^{16 \times (d+1)}, \quad \mathbf{x} \triangleq \begin{bmatrix} x_1 \\ \vdots \\ x_{d+1} \end{bmatrix} \in \mathbb{R}^{d+1}.$$

The solution $\hat{\mathbf{x}} = \mathbf{A}^\dagger \mathbf{y}$ yields the the desired optimal least-squares estimate of the coefficients of the degree- d polynomial $p_d(t) = \sum_{i=1}^{d+1} x_i t^{i-1}$ that minimizes the error $\|\mathbf{y} - \mathbf{A}\mathbf{x}\|_2$. The \mathbf{A} of the form above is called a Vandermonde matrix.

Here is **Julia** code:

```
T = [2.1e-3, 0.136, 0.268, 0.402, 0.536,
      0.668, 0.802, 0.936, 1.068, 1.202,
      1.336, 1.468, 1.602, 1.736, 1.868, 2.000]

#f = 0.5 * exp.(0.8*T) # part a
f = 0.5 * exp.(0.8*T) + randn(length(T)) # part b
d = length(T)-1

A15 = [T.^j for TT in T, j=0:d]
A2 = A15[:,1:3]
x16 = pinv(A15)*f
x2 = pinv(A2)*f
t = linspace(0,2,500)


using Plots; pyplot()
plot(t, 0.5*exp.(0.8*t), color=:red, label = "f(t)", ylim=(-1,4))
scatter!(T,f, marker=:circle, label = "Samples")

a16 = [t.^j for tt in t, j=0:d]
a2 = a16[:,1:3]


plot!(t, (a16*x16), color=:black, label = "d = 15 fit")
plot!(t, (a2*x2), color=:blue, label = "d = 2 fit", xaxis = "t", yaxis = "Value")
savefig("tmp2.pdf")
```

This yielded the upper figure below. Note that the error for $d = 15$ is much smaller than for $d = 2$. In fact the $d = 15$ error is exactly zero (to within numerical precision) because $\text{rank}(\mathbf{A}) = 16 = \dim(\mathbf{y})$ so the solution is exact! Comparing the least-squares coefficients obtained for $d = 2$ and $d = 15$ reveals that the first three terms are about the same. When noise is added we get the lower figure below. The polynomial for $d = 15$ passes *exactly* through the noisy samples. This property follows from the previous argument because $\text{rank}(\mathbf{A}) = 16$ so that error will be zero, relative to the samples, but large relative to the function. This example illustrates why choosing large model orders ($= d$) can make things worse by “over fitting” noise.

Part (a) - when there is no noise:



Part (b) - when there is noise.



Use the `cond(A)` command in **Julia** to obtain the condition number: the ratio of the largest and smallest singular value of a matrix. The condition number is important because the measurement error manifests as a $\delta \mathbf{y}$ so that the error in the least squares solution (relative to the noise-less case) is given by exactly by $\delta \hat{\mathbf{x}} = \mathbf{A}^\dagger(\mathbf{y} + \delta \mathbf{y}) - \mathbf{A}^\dagger \mathbf{y} = \mathbf{A}^\dagger \delta \mathbf{y}$. A large condition number means that small $\delta \mathbf{y}$ can get amplified, producing large $\delta \mathbf{x}$ and hence instabilities as manifested in the overfitting.

An estimate of the degree to which $\delta \mathbf{x}$ can change as a function of the condition number κ is given by the relation:

$$\frac{\|\delta \mathbf{x}\|}{\|\mathbf{x}\|} \leq \kappa \frac{\|\delta \mathbf{y}\|}{\|\mathbf{y}\|}.$$

Here $\|\delta \mathbf{y}\|$ is the norm of the noise. The condition number $\kappa = \sigma_1(\mathbf{A})/\sigma_{\min(m,n)}(\mathbf{A})$ is large when $d = 15$ (check it) compared to when $d = 2$.

Pr. 8.

Here $\mathbf{A} = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$ and $\mathbf{A}^\dagger = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$ and the optimal least-squares estimate is: $\mathbf{x}^* = \mathbf{A}^\dagger \mathbf{b} = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$.

- (a) Consider $\mathbf{A}_1 = \begin{bmatrix} 1 & \delta \\ 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \begin{bmatrix} 1 & \delta \end{bmatrix}$ so that by inspection an SVD of \mathbf{A}_1 is $\mathbf{A}_1 = \sigma_1 \mathbf{u}_1 \mathbf{v}_1^T$ where $\sigma_1 = \sqrt{1 + \delta^2}$, $\mathbf{u}_1 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$ and $\mathbf{v}_1 = \begin{bmatrix} 1/\sqrt{1 + \delta^2} \\ \delta/\sqrt{1 + \delta^2} \end{bmatrix}$. We do not need to compute the other singular vectors because the optimal (minimum-norm) least-squares solution is simply: $\mathbf{z}^* = \mathbf{A}_1^\dagger \mathbf{b} = \frac{1}{\sigma_1} \mathbf{v}_1 \mathbf{u}_1^T \mathbf{b} = \frac{1}{1 + \delta^2} \begin{bmatrix} 1 \\ \delta \end{bmatrix}$, for which $\mathbf{z}^* \rightarrow \begin{bmatrix} 1 \\ 0 \end{bmatrix}$ as $\delta \rightarrow 0$, so $\|\mathbf{z}^* - \mathbf{x}^*\|_2 \rightarrow 0$.
- (b) Consider now the case where $\mathbf{A}_2 = \begin{bmatrix} 1 & 0 \\ 0 & \delta \end{bmatrix}$, which has $\text{rank}(\mathbf{A}) = 2$ for $\delta > 0$. Here the optimal least-squares solution is: $\mathbf{z}^* = \mathbf{A}_2^\dagger \mathbf{b} = \begin{bmatrix} 1 \\ 1/\delta \end{bmatrix}$ so that $\|\mathbf{x}^* - \mathbf{z}^*\|_2 \rightarrow \infty$ as $\delta \rightarrow 0$.

This exercise illustrates that the manner in which error is introduced can impact a solution negligibly or dramatically!

Pr. 9.

Let $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^H$. If $\text{rank}(\mathbf{A}) = n$, then \mathbf{A} has n linearly independent columns, $n \leq m$, and $\mathbf{A}^H \mathbf{A}$ is invertible. Then:

$$\begin{aligned}
 (\mathbf{A}^H \mathbf{A})^{-1} \mathbf{A}^H &= (\mathbf{V} \mathbf{\Sigma}^H \mathbf{U}^H \mathbf{U} \mathbf{\Sigma} \mathbf{V}^H)^{-1} \mathbf{V} \mathbf{\Sigma}^H \mathbf{U}^H = (\mathbf{V} \mathbf{\Sigma}^H \mathbf{\Sigma} \mathbf{V}^H)^{-1} \mathbf{V} \mathbf{\Sigma}^H \mathbf{U}^H \\
 &= \mathbf{V} (\mathbf{\Sigma}^H \mathbf{\Sigma})^{-1} \mathbf{\Sigma}^H \mathbf{U}^H = \mathbf{V} \begin{bmatrix} \frac{1}{\sigma_1^2} & 0 & \dots & 0 \\ 0 & \frac{1}{\sigma_2^2} & \dots & 0 \\ & \ddots & \ddots & \\ 0 & 0 & \dots & \frac{1}{\sigma_n^2} \end{bmatrix} \begin{bmatrix} \sigma_1 & 0 & \dots & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 & \dots & 0 \\ & \ddots & & & \ddots & \\ 0 & 0 & \dots & \sigma_n & \dots & 0 \end{bmatrix} \mathbf{U}^H \\
 &= \mathbf{V} \begin{bmatrix} \frac{1}{\sigma_1} & 0 & \dots & 0 & \dots & 0 \\ 0 & \frac{1}{\sigma_2} & \dots & 0 & \dots & 0 \\ & \ddots & & & \ddots & \\ 0 & 0 & \dots & \frac{1}{\sigma_n} & \dots & 0 \end{bmatrix} \mathbf{U}^H = \mathbf{A}^\dagger.
 \end{aligned}$$

Pr. 10.

A possible Julia implementation is

```

function compute_normals(I, L)
#
# Syntax:      N = compute_normals(I, L)
#
# Inputs:      I is an m x n x d matrix whose d slices contain m x n images
#              of a common scene under different lighting conditions
#
#              L is a 3 x d matrix whose columns are the lighting direction
#              vectors for the images in I, with d >= 3
#
# Outputs:     N is an m x n x 3 matrix containing the unit-norm
#              surface normal vectors for each pixel in the scene
#
#
m, n, d = size(I) # Parse inputs
L = mapslices(normalize, L, 1) # Normalize lighting direction vectors

# Solve least squares problem
# Here, using pinv() is efficient because L is a small matrix
# and we apply pinv(L) to many (mn) pixels.
I = reshape(I, m * n, d)
N = I * pinv(L)
N = reshape(N, m, n, 3)

# alternative one-line "Julia way" that avoids reshape:
# N = mapslices((v) -> pinv(L')*v, I, 3) # apply pinv to each pixel
# However, it ran slower than the above, per autograder output.

N = mapslices(normalize, N, 3) # Normalize normal vectors
return N
end

```

