

# Lecture 11: Deep RL (RL Algorithms based on Neural Networks)

Course: Reinforcement Learning Theory  
Instructor: Lei Ying  
Department of EECS  
University of Michigan, Ann Arbor

- Motivation: Q-learning overestimates action values under certain conditions, because of the maximization step over estimated action values.
- Question: How bad is overestimation?

# Deep Q-learning

- Thrun and Schwartz (1993):

Assume under function approximation,

$$\tilde{Q}(s, a) = Q(s, a) + Y_{s,a}$$

where  $E[Y_{s,a}] = 0$  (zero mean).

Q-learning: with transition  $(s, a, s')$ ,

$$Q(s, a) \leftarrow r(s, a) + \alpha \max_{\hat{a}} Q(s', \hat{a})$$

# Deep Q-learning

Error from noise:

$$\begin{aligned} Z &= r(s, a) + \alpha \max_{\hat{a}} \tilde{Q}(s', \hat{a}) - (r(s, a) + \alpha \max_{\hat{a}} Q(s', \hat{a})) \\ &= \alpha \left( \max_{\hat{a}} \tilde{Q}(s', \hat{a}) - \max_{\hat{a}} Q(s', \hat{a}) \right) \\ &= -\alpha \left( \max_{\hat{a}} Q(s', \hat{a}) - \max_{\hat{a}} (Q(s', \hat{a}) + Y_{s', \hat{a}}) \right) \end{aligned}$$

Claim:  $E[Y_{s,a}] = 0 \xrightarrow{\text{often}} E[Z] > 0$

- Consider the case  $Q(s', \hat{a}_i) = Q(s', \hat{a}_j) \quad \forall i \neq j$   
i.e. Q-values are the same for all actions at state  $s'$ .
- Assume  $Y_{s', \hat{a}} \sim \text{uniform}[-\epsilon, \epsilon]$  and independent across  $\hat{a}$

$$\begin{aligned} E[Z] &= \alpha E[-Q(s', \hat{a}_1) + Q(s', \hat{a}_1) + \max_{\hat{a}} Y_{s', \hat{a}}] \\ &= \alpha E[\max_{\hat{a}} Y_{s', \hat{a}}] \end{aligned}$$

Let  $n$  be the number of actions,

$$\begin{aligned} &= \alpha \int_{-\infty}^{\infty} x n f(x) \underbrace{\left( \int_{-\infty}^x f(z) dz \right)}_{P(\leq x)}^{n-1} dx \\ &= \alpha n \int_{-\epsilon}^{\epsilon} x \frac{1}{2\epsilon} \left( \frac{1}{2} + \frac{x}{2\epsilon} \right)^{n-1} dx \end{aligned}$$

- Define  $y = \frac{1}{2} + \frac{x}{2\epsilon}$ , we have

$$\begin{aligned} E[Err] &= an \int_0^1 (2\epsilon y - \epsilon) y^{n-1} dy \\ &= \alpha n \epsilon \int_0^1 2y^n - y^{n-1} dy \\ &= \alpha n \epsilon \left( \frac{2}{n+1} - \frac{1}{n} \right) \\ &= \alpha \epsilon \frac{n-1}{n+1} \end{aligned}$$

- Remark: This is the worst case. The error is smaller if  $Q(s', \hat{a}_1) \neq Q(s', \hat{a}_2)$

# Deep Q-learning

- Lower bound (van Hasselt, Guez, Silver (2015)): again, consider  $Q(s', \hat{a}_1) = Q(s', \hat{a}_2) \quad \forall \hat{a}_1 \neq \hat{a}_2$
- Define  $Q(s', \hat{a}) = V_*(s') \quad \forall \hat{a}$ , and assume

$$\sum_{\hat{a}} (\tilde{Q}(s', \hat{a}) - V_*(s')) = 0 \text{ (unbiased as a whole)}$$

$$\text{and } \frac{1}{n} \sum_{\hat{a}} (\tilde{Q}(s', \hat{a}) - V_*(s'))^2 = c$$

Claim:

$$\max_{\hat{a}} \tilde{Q}(s', \hat{a}) \geq V_*(s') + \sqrt{\frac{c}{n-1}}$$

# Deep Q-learning

Proof (by contradiction):

- Define  $\epsilon_{\hat{a}} = \tilde{Q}(s', \hat{a}) - V_*(s')$

Suppose there exist  $\{\epsilon_{\hat{a}}\}$  such that  $\epsilon_{\hat{a}} < \sqrt{\frac{c}{n-1}} \quad \forall \hat{a}$

- Let  $\{\epsilon_i^+\}_{i=1,\dots,m}$  be positive  $\epsilon$ , and  $\{\epsilon_j^-\}_{j=1,\dots,n-m}$  be negative  $\epsilon$ .

$$\begin{aligned} \sum_{j=1}^{n-m} |\epsilon_j^-| &= \sum_{i=1}^m \epsilon_i^+ \quad (\text{because } \sum_{i=1}^m \epsilon_i = 0) \\ &\leq m \max_{i=1,\dots,m} \epsilon_i^+ < m \sqrt{\frac{c}{n-1}} \\ \implies \sum_{j=1}^{n-m} (\epsilon_j^-)^2 &\leq \left( \sum_{j=1}^{n-m} |\epsilon_j^-| \right)^2 \\ &< m^2 \frac{c}{n-1} \end{aligned}$$



$$\begin{aligned}\Rightarrow \sum_{i=1}^n \epsilon_i^2 &= \sum_{i=1}^m (\epsilon_i^+)^2 + \sum_{j=1}^{n-m} (\epsilon_j^-)^2 \\ &< m \frac{c}{n-1} + m^2 \frac{c}{n-1} \\ &= \frac{m(m+1)}{n-1} c \\ &\leq nc \text{ (because } m \leq n-1 \text{)}.\end{aligned}$$

In other words,

$$\sum_{i=1}^n \epsilon_i^2 < nc \text{ (because } m \leq n-1 \text{). Contradiction}$$

# Deep Q-learning

- Overestimation leads to the failure of Q-learning
- Example (Thrun and Schwartz (1993)):
  1. A set of goal states
  2.  $r_s = \begin{cases} 1 & s \text{ is a goal state} \\ 0 & \text{otherwise} \end{cases}$
  3. state-transition is deterministic
- Suppose  $\langle s_i, a_i \rangle (i \in \{0, \dots, L\})$  is an optimal state-action sequence.
- Then, necessary condition:  $Q(s_i, a_i) < Q(s_{i+1}, a_{i+1})$  when  $\alpha < 1$  (because the reward is only received at the end).

# Deep Q-learning

- Define  $c = \epsilon \frac{n-1}{n+1}$  with the same uniform noise.
- Suppose the algorithm overestimates  $Q$  values by  $\alpha c$ ,  $c = \epsilon \frac{n-1}{n+1}$  every step. Then,

$$q_L = 1$$

$$q_{L-1} = \alpha + \alpha c = \alpha(1 + c)$$

$$q_{L-2} = \underbrace{\alpha q_{L-1}}_{\text{true estimation}} + \underbrace{\alpha c}_{\text{error}}$$

$$= \alpha(q_{L-1} + c)$$

$$= \alpha(\alpha + \alpha c + c)$$

$$= \alpha^2 + \alpha^2 c + \alpha c$$

$$q_i = \alpha^{L-i} + \sum_{k=1}^{L-i} \alpha^k c$$

- Necessary condition:

$$\begin{aligned} 0 \geq q_i - q_{i+1} &= \alpha^{L-i} + \sum_{k=1}^{L-i} \alpha^k c - \alpha^{L-i-1} - \sum_{k=1}^{L-i-1} \alpha^k c \\ &= \alpha^{L-i-1}(\alpha - 1) + \alpha^{L-i} c \end{aligned}$$

$$0 \geq \alpha - 1 + \alpha c \iff \alpha \leq \frac{1}{1+c}$$

Otherwise, Q-learning may fail.

# Q-learning methods overview

- Q-learning (Watkin '89) tabular method.

$$Q_{k+1}(i, u) = Q_k(i, u) + \beta_k(r(i, u) + \alpha \max_v Q_k(j, v) - Q_k(i, u))$$

- Q-learning (function approximation)

$$\theta_{k+1} = \theta_k + \beta_k \left( r(s_t, a_t) + \alpha \max_u Q(s_{t+1}, u; \theta_t) - Q(s_t, a_t; \theta_t) \right) \times \nabla_{\theta_t} Q(s_t, a_t; \theta_t)$$

# Q-learning methods overview

- Double Q-learning (van Hasselt 2010)

Maintain two Q-functions, denoted by their parameters  $\theta$  and  $\theta'$ ,

$$\theta_{k+1} = \theta_k + \beta_k \delta_k \nabla_{\theta_k} Q(s_k, a_k; \theta_k),$$

where

$$\delta_k = r(s_k, a_k) + \alpha Q(s_{k+1}, \arg \max_u Q(s_{k+1}, u; \theta_k); \theta'_k) - Q(s_k, a_k; \theta_k),$$

and

$$\theta'_{k+1} = \theta'_k + \beta_k \delta'_k \nabla_{\theta'_k} Q(s_k, a_k; \theta'_k)$$

where

$$\delta'_k = r(s_k, a_k) + \alpha Q(s_{k+1}, \arg \max_u Q(s_{k+1}, u; \theta'_k); \theta_k) - Q(s_k, a_k; \theta'_k).$$

# The double estimator to estimate $\max_i \mathbb{E}[X_i]$

(Van Hasselt '2010)

$S_i$ : set of samples from  $X_i$

**Single estimator:**  $\max_i \mathbb{E}[X_i] \approx \max_i \frac{1}{|S_i|} \sum_{s \in S_i} s$

**Double estimator:**

Split  $S_i$  into  $S_i^A$  and  $S_i^B$  such that  $S_i^A \cup S_i^B$  and  $S_i^A \cap S_i^B = \emptyset$

Define

$$\mu_i^A = \frac{1}{|S_i^A|} \sum_{s \in S_i^A} s, \quad \mu_i^B = \frac{1}{|S_i^B|} \sum_{s \in S_i^B} s$$

- pick  $i^*$  such that  $\mu_{i^*}^A = \max_i \mu_i^A$
- approximate  $\max_i \mathbb{E}[X_i] = \mu_{i^*}^B$

**Remark:** Separate the maximizer and evaluation.

**Claim:** Double estimator is an underestimator.

# The double estimator to estimate $\max_i \mathbb{E}[X_i]$

Proof.

If  $a^*$  is the maximizer,

$$\mathbb{E}[\mu_{a^*}^B] = \mathbb{E}[X_{a^*}] = \max_i \mathbb{E}[X_i].$$

If not,  $\mathbb{E}[\mu_{a^*}^B] < \max_i \mathbb{E}[X_i]$ .

$$\Rightarrow \mathbb{E}[\mu_{a^*}^B] = \mathbb{P}(a^* \in M) \mathbb{E}[\mu_{a^*}^B | a^* \in M] + \mathbb{P}(a^* \notin M) \mathbb{E}[\mu_{a^*}^B | a^* \notin M] \quad (1)$$

$$< \mathbb{P}(a^* \in M) \max_i \mathbb{E}[X_i] + \mathbb{P}(a^* \notin M) \max_i \mathbb{E}[X_i] \quad (2)$$

$$= \max_i \mathbb{E}[X_i] \quad (3)$$

where  $M$  is the set of maximizers.

If  $\{X_i\}$  are i.i.d., then  $\mathbb{E}[\mu_{a^*}^B] = \max_i \mathbb{E}[X_i]$ .





# Q-learning methods overview

- Deep Q-learning (DQN) (Mnih et al. 2015)

Maintain a target network  $\theta^-$ ,

$\theta_k^- \leftarrow \theta_k$  every  $\tau$  steps.

$$Q_{k+1} = \theta_k + \beta_t(Y_t - Q(s_k, a_t; \theta_k))\nabla_{\theta_k} Q(s_k, a_k; \theta_k)$$

$$Y_k = r(s_k, a_k) + \alpha \max_u Q(s_{k+1}, u; Q_k^-)$$

target network

Action selection and evaluation based on the target network.

# Q-learning methods overview

- Double DQN (van Hasselt, Guez, Silver 2015)

$$Y_k = r(s_k, a_k) + \alpha Q(s_{k+1}, \arg \max_u Q(s_{k+1}, u; Q_k); \theta_k^-)$$

action selection based on online networks

evaluation based on target network

# Q-learning methods overview

- Clipped Double-Q Learning (Fujimoto, van Hoof, David Meger 2018)  
Maintain two Q-functions, denoted by their parameters  $\theta$  and  $\theta'$ ,

$$\begin{aligned}\theta_{k+1} &= \theta_k + \beta_k (Y_k - Q(s_t, a_t; \theta_t)) \nabla_{\theta_k} Q(s_k, a_k; \theta_k) \\ \theta'_{k+1} &= \theta'_k + \beta_k (Y_k - Q(s_t, a_t; \theta'_t)) \nabla_{\theta'_k} Q(s_k, a_k; \theta'_k),\end{aligned}$$

where

$$Y_k = r(s_k, a_k) + \alpha \min \left\{ \max_u Q(s_k + 1, u; \theta_k), \max_u Q(s_k + 1, u; \theta'_k) \right\}$$

# References

- Sebastian Thrun and A. Schwartz, *Issues in Using Function Approximation for Reinforcement Learning*, Proceedings of the Connectionist Models Summer School, June, 1993.
- H Hasselt, *Double Q-learning*, NeurIPS, 2010.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, Martin Riedmiller, *Playing atari with deep reinforcement learning*, arXi, 2013
- H. van Hasselt. , A. Guez, D. Silver, *Deep Reinforcement Learning with Double Q-Learning*, AAAI. 2016.
- Scott Fujimoto, Herke van Hoof, and David Meger, *Addressing Function Approximation Error in Actor-Critic Methods*, ICML, 2018.

---

**Acknowledgements:** I would like to thank Alex Zhao for helping prepare the slides, and Honghao Wei and Zixian Yang for correcting typos/mistakes.