

Pr. 1.(a) $\tilde{x} = \beta_1/\beta_3$, $\tilde{y} = \beta_2/\beta_3$.(b) From the first and third elements of vector β , we get

(1)
$$\beta_3 \tilde{x} = h_1^T \alpha = \alpha^T h_1$$

(2)
$$\beta_3 = h_3^T \alpha = \alpha^T h_3$$

from which we see that

$$\alpha^T h_1 - \tilde{x} \alpha^T h_3 = 0.$$

In matrix-vector form:

$$\begin{bmatrix} \alpha^T & \mathbf{0}^T & -\tilde{x} \alpha^T \end{bmatrix} \begin{bmatrix} h_1 \\ h_2 \\ h_3 \end{bmatrix} = \begin{bmatrix} \alpha^T & \mathbf{0}^T & -\tilde{x} \alpha^T \end{bmatrix} \text{vec}(\mathbf{H}) = 0.$$

Following the same argument with the second and third elements of β , we get

(3)
$$\beta_3 \tilde{y} = h_2^T \alpha = \alpha^T h_2$$

$$\beta_3 = h_3^T \alpha = \alpha^T h_3$$

from which

$$\begin{bmatrix} \mathbf{0}^T & \alpha^T & -\tilde{y} \alpha^T \end{bmatrix} \begin{bmatrix} h_1 \\ h_2 \\ h_3 \end{bmatrix} = 0$$

so that the required matrix A is

$$A = \begin{bmatrix} \alpha^T & \mathbf{0}^T & -\tilde{x} \alpha^T \\ \mathbf{0}^T & \alpha^T & -\tilde{y} \alpha^T \end{bmatrix}.$$

Alternatively, one could combine (1), (2) and (3) to see that

$$\alpha^T h_1 + \alpha^T h_2 - (\tilde{x} + \tilde{y}) h_3 = \beta_3 \tilde{x} + \beta_3 \tilde{y} - (\tilde{x} + \tilde{y}) \beta_3 = 0,$$

so that

$$A = \begin{bmatrix} \alpha^T & \alpha^T & -(\tilde{x} + \tilde{y}) \alpha^T \\ \alpha^T & \alpha^T & -(\tilde{x} + \tilde{y}) \alpha^T \end{bmatrix}.$$

(c) Vector \mathbf{h} is in the *null space* of A , i.e., $\mathbf{h} \in \mathcal{N}(A)$.**Pr. 2.**

$$\begin{aligned} (\mathbf{A}^\dagger \mathbf{b})^T (\mathbf{I} - \mathbf{A}^\dagger \mathbf{A}) \mathbf{y} &= \mathbf{b}^T (\mathbf{A}^\dagger)^T (\mathbf{I} - \mathbf{V} \Sigma^\dagger \mathbf{U}^T \mathbf{U} \Sigma \mathbf{V}^T) \mathbf{y} = \mathbf{b}^T \mathbf{U} (\Sigma^\dagger)^T \mathbf{V}^T (\mathbf{V} \mathbf{V}^T - \mathbf{V} \Sigma^\dagger \Sigma \mathbf{V}^T) \mathbf{y} \\ &= \mathbf{b}^T \mathbf{U} ((\Sigma^\dagger)^T - (\Sigma^\dagger)^T \Sigma^\dagger \Sigma) \mathbf{V}^T \mathbf{y}. \end{aligned}$$

If $\Sigma \in \mathbb{R}^{m \times n} = \text{diag}(\sigma_1, \sigma_2, \dots, \sigma_r)$, then we can verify that $(\Sigma^\dagger)^T = (\Sigma^\dagger)^T \Sigma^\dagger \Sigma$ through direct multiplication, so that $(\mathbf{A}^\dagger \mathbf{b})^T (\mathbf{I} - \mathbf{A}^\dagger \mathbf{A}) \mathbf{y} = 0$.**Pr. 3.**(a) Every point $[x \ y \ z]^T$ on the plane satisfies the equation $\begin{bmatrix} a & b & c \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix} = 0$. Therefore every point on the plane

$$\text{must satisfy } \begin{bmatrix} x \\ y \\ z \end{bmatrix} \in \mathcal{N}(\begin{bmatrix} a & b & c \end{bmatrix}). \text{ Clearly: } \begin{bmatrix} a & b & c \end{bmatrix} = \underbrace{1}_{=u_1} \underbrace{\sqrt{a^2 + b^2 + c^2}}_{=\sigma_1} \underbrace{\frac{\begin{bmatrix} a & b & c \end{bmatrix}}{\sqrt{a^2 + b^2 + c^2}}}_{=\mathbf{v}_1^T}.$$

Note that $\mathcal{N}(\begin{bmatrix} a & b & c \end{bmatrix}) = \text{span}(\{\mathbf{v}_2, \mathbf{v}_3\})$, since we are considering a rank-1 matrix. Therefore $\{\mathbf{v}_2, \mathbf{v}_3\}$ are an orthonormal basis for the plane. Two basis vectors are required to express every point on the plane.

(b) The nearest point on the plane is given by

$$P_{\mathcal{R}\{\mathbf{v}_2, \mathbf{v}_3\}} \begin{bmatrix} \alpha \\ \beta \\ \gamma \end{bmatrix} = (\mathbf{v}_2 \mathbf{v}_2^T + \mathbf{v}_3 \mathbf{v}_3^T) \begin{bmatrix} \alpha \\ \beta \\ \gamma \end{bmatrix} = (\mathbf{I} - \mathbf{v}_1 \mathbf{v}_1^T) \begin{bmatrix} \alpha \\ \beta \\ \gamma \end{bmatrix}.$$

Pr. 4.

Define: $\mathbf{A} = [\mathbf{q}_1 \quad \mathbf{q}_2]$, and $\mathbf{x} = \begin{bmatrix} \alpha \\ \beta \end{bmatrix}$.

(a) The linear least squares estimate that minimizes $\|\mathbf{Ax} - \mathbf{b}\|_2$ is given by $\hat{\mathbf{x}} = \mathbf{A}^\dagger \mathbf{b}$. The SVD of \mathbf{A} is simply:

$$\mathbf{A} = \sum_{i=1}^2 1 \mathbf{q}_i \mathbf{e}_i^T,$$

where \mathbf{e}_i denotes the i th unit vector. Thus:

$$\mathbf{A}^\dagger = \sum_{i=1}^2 1 \mathbf{e}_i \mathbf{q}_i^T,$$

and hence

$$\hat{\mathbf{x}} = \mathbf{A}^\dagger \mathbf{b} = \sum_{i=1}^2 1 \mathbf{e}_i \mathbf{q}_i^T \mathbf{b} = \mathbf{A}^T \mathbf{b}.$$

This solution is unsurprising, because what we get is precisely the first two “coordinates” of the vector \mathbf{b} relative to the basis whose first two basis vectors correspond to the columns of \mathbf{A} .

(b) Here we have that the residual (or error) vector:

$$\mathbf{r} = \mathbf{b} - \mathbf{A}\hat{\mathbf{x}} = (\mathbf{I} - \mathbf{A}\mathbf{A}^T)\mathbf{b} = \sum_{i=3}^n \mathbf{q}_i \mathbf{q}_i^T \mathbf{b},$$

where \mathbf{q}_i for $i = 3, \dots, n$ are the $n-2$ (unit norm) basis vectors, orthogonal to \mathbf{q}_1 and \mathbf{q}_2 so that $\text{span}(\{\mathbf{q}_1, \dots, \mathbf{q}_n\}) = \mathbb{R}^n$. Hence $\mathbf{q}_1^T \mathbf{r} = \sum_{i=3}^n \mathbf{q}_1^T \mathbf{q}_i \mathbf{b} = 0$ and $\mathbf{q}_2^T \mathbf{r} = \sum_{i=3}^n \mathbf{q}_2^T \mathbf{q}_i \mathbf{b} = 0$.

Pr. 5.

Here

$$\mathbf{A} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} = \underbrace{\frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}}_{=\mathbf{U}} \underbrace{\begin{bmatrix} 2 & 0 \\ 0 & 0 \end{bmatrix}}_{=\mathbf{\Sigma}} \underbrace{\frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}^T}_{=\mathbf{V}^T},$$

since \mathbf{A} is rank-1 and can be written as an outerproduct $\mathbf{A} = \mathbf{z}\mathbf{z}^T$ where $\mathbf{z} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$ and we know the eigen-decomposition of such rank-1 matrices from previous homeworks. Recall that since \mathbf{A} is symmetric, positive-semidefinite, its eigen-decomposition is the same as a singular value decomposition.

Consider the minimum norm solution given by

$$\begin{aligned} \hat{\mathbf{x}} = \mathbf{A}^\dagger \mathbf{b} &= \begin{bmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ 1/\sqrt{2} & -1/\sqrt{2} \end{bmatrix} \begin{bmatrix} 1/2 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ 1/\sqrt{2} & -1/\sqrt{2} \end{bmatrix}^T \mathbf{b} \\ &= \begin{bmatrix} 0.25 & 0.25 \\ 0.25 & 0.25 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 0.75 \\ 0.75 \end{bmatrix}. \end{aligned}$$

Here $\text{rank}(\mathbf{A}) = 1 < 2$. Moreover

$$\ker(\mathbf{A}) = \text{span} \left(\begin{bmatrix} 1 \\ -1 \end{bmatrix} \right),$$

so that for $\alpha \in \mathbb{R}$, all

$$\tilde{\mathbf{x}} = \hat{\mathbf{x}} + \alpha \begin{bmatrix} 1 \\ -1 \end{bmatrix}$$

will yield the same the (minimum) squared error.

Pr. 6.

(a)

$$\begin{aligned}
& (U_x \otimes U_y)(\Sigma_x \otimes \Sigma_y)(V_x \otimes V_y)^H \\
&= (U_x \otimes U_y)(\Sigma_x \otimes \Sigma_y)(V_x^H \otimes V_y^H) \\
&= (U_x \otimes U_y)((\Sigma_x V_x^H) \otimes (\Sigma_y V_y^H)) \\
&= (U_x \Sigma_x V_x^H) \otimes (U_y \Sigma_y V_y^H) = X \otimes Y.
\end{aligned}$$

To show that $U_x \otimes U_y$ is unitary, observe that

$$\begin{aligned}
& (U_x \otimes U_y)^H (U_x \otimes U_y) \\
&= (U_x^H \otimes U_y^H)(U_x \otimes U_y) \\
&= (U_x^H U_x) \otimes (U_y^H U_y) \\
&= I_m \otimes I_p = I_{mp}.
\end{aligned}$$

The above argument can be repeated to show that

- $(U_x \otimes U_y)(U_x \otimes U_y)^H = I_{mp}$
- $(V_x \otimes V_y)^H (V_x \otimes V_y) = I_{nq}$
- $(V_x \otimes V_y)(V_x \otimes V_y)^H = I_{nq}$.

(b)

$$\begin{aligned}
& (Q_a \otimes Q_b)(\Lambda_a \otimes \Lambda_b)(Q_a \otimes Q_b)^H \\
&= (Q_a \otimes Q_b)(\Lambda_a \otimes \Lambda_b)(Q_a^H \otimes Q_b^H) \\
&= (Q_a \otimes Q_b)(\Lambda_a Q_a^H \otimes \Lambda_b Q_b^H) \\
&= Q_a \Lambda_a Q_a^H \otimes Q_b \Lambda_b Q_b^H = A \otimes B.
\end{aligned}$$

To show that $Q_a \otimes Q_b$ is unitary, observe that

$$\begin{aligned}
& (Q_a \otimes Q_b)^H (Q_a \otimes Q_b) \\
&= (Q_a^H \otimes Q_b^H)(Q_a \otimes Q_b) \\
&= (Q_a^H Q_a \otimes Q_b^H Q_b) \\
&= I_n \otimes I_m = I_{mn}.
\end{aligned}$$

The above argument can be repeated to show that $(Q_a \otimes Q_b)(Q_a \otimes Q_b)^H = I_{mn}$.

(c)

$$\begin{aligned}
A \oplus B &= A \otimes I_m + I_n \otimes B \\
&= A \otimes (Q_b I_m Q_b^H) + (Q_a I_n Q_a^H) \otimes B \\
\text{by (b)} \rightarrow &= (Q_a \otimes Q_b)(\Lambda_a \otimes I_m)(Q_a \otimes Q_b)^H + (Q_a \otimes Q_b)(I_n \otimes \Lambda_b)(Q_a \otimes Q_b)^H \\
&= (Q_a \otimes Q_b)(\Lambda_a \otimes I_m + I_n \otimes \Lambda_b)(Q_a \otimes Q_b)^H \\
&= (Q_a \otimes Q_b)(\Lambda_a \oplus \Lambda_b)(Q_a \otimes Q_b)^H.
\end{aligned}$$

As in (b), $Q_a \otimes Q_b$ is unitary.

Pr. 7.

We have to show that $\|A(\lambda x + (1-\lambda)y) - b\|_2 \leq \lambda \|Ax - b\|_2 + (1-\lambda) \|Ay - b\|_2$.

To that end note that

$$\begin{aligned}
f(\lambda x + (1-\lambda)y) &= \|A(\lambda x + (1-\lambda)y) - b\|_2 \\
&= \|A(\lambda x + (1-\lambda)y) - b(\lambda + 1 - \lambda)\|_2 \\
&= \|(\lambda Ax - \lambda b) + ((1-\lambda)Ay - (1-\lambda)b)\|_2 \\
&\leq \|\lambda(Ax - b)\|_2 + \|(1-\lambda)(Ay - b)\|_2, \quad (\text{via the triangle inequality}) \\
&\leq \lambda \|Ax - b\|_2 + (1-\lambda) \|Ay - b\|_2 \\
&\leq \lambda f(x) + (1-\lambda)f(y),
\end{aligned}$$

establishing convexity. Next we need to show that, given two matrices A, B , $\forall \lambda \in [0, 1]$:

$$\sigma_1(\lambda A + (1-\lambda)B) \leq \lambda \sigma_1(A) + (1-\lambda) \sigma_1(B).$$

From the hint:

$$\sigma_1(\lambda \mathbf{A} + (1 - \lambda) \mathbf{B}) = \max_{\|\mathbf{u}\|_2=1} \|(\lambda \mathbf{A} + (1 - \lambda) \mathbf{B}) \mathbf{u}\|_2$$

By triangle inequality of norms,

$$\begin{aligned} &\leq \max_{\|\mathbf{u}\|_2=1} \|\lambda \mathbf{A} \mathbf{u}\|_2 + \|(1 - \lambda) \mathbf{B} \mathbf{u}\|_2 \\ &\leq \max_{\|\mathbf{u}\|_2=1} \lambda \|\mathbf{A} \mathbf{u}\|_2 + (1 - \lambda) \|\mathbf{B} \mathbf{u}\|_2 \\ &\leq \max_{\|\mathbf{u}\|_2=1} \lambda \|\mathbf{A} \mathbf{u}\|_2 + \max_{\|\mathbf{u}\|_2=1} (1 - \lambda) \|\mathbf{B} \mathbf{u}\|_2 \\ &= \lambda \sigma_1(\mathbf{A}) + (1 - \lambda) \sigma_1(\mathbf{B}). \end{aligned}$$

Thus, $\sigma_1(\cdot)$ is a convex function.

Pr. 8.

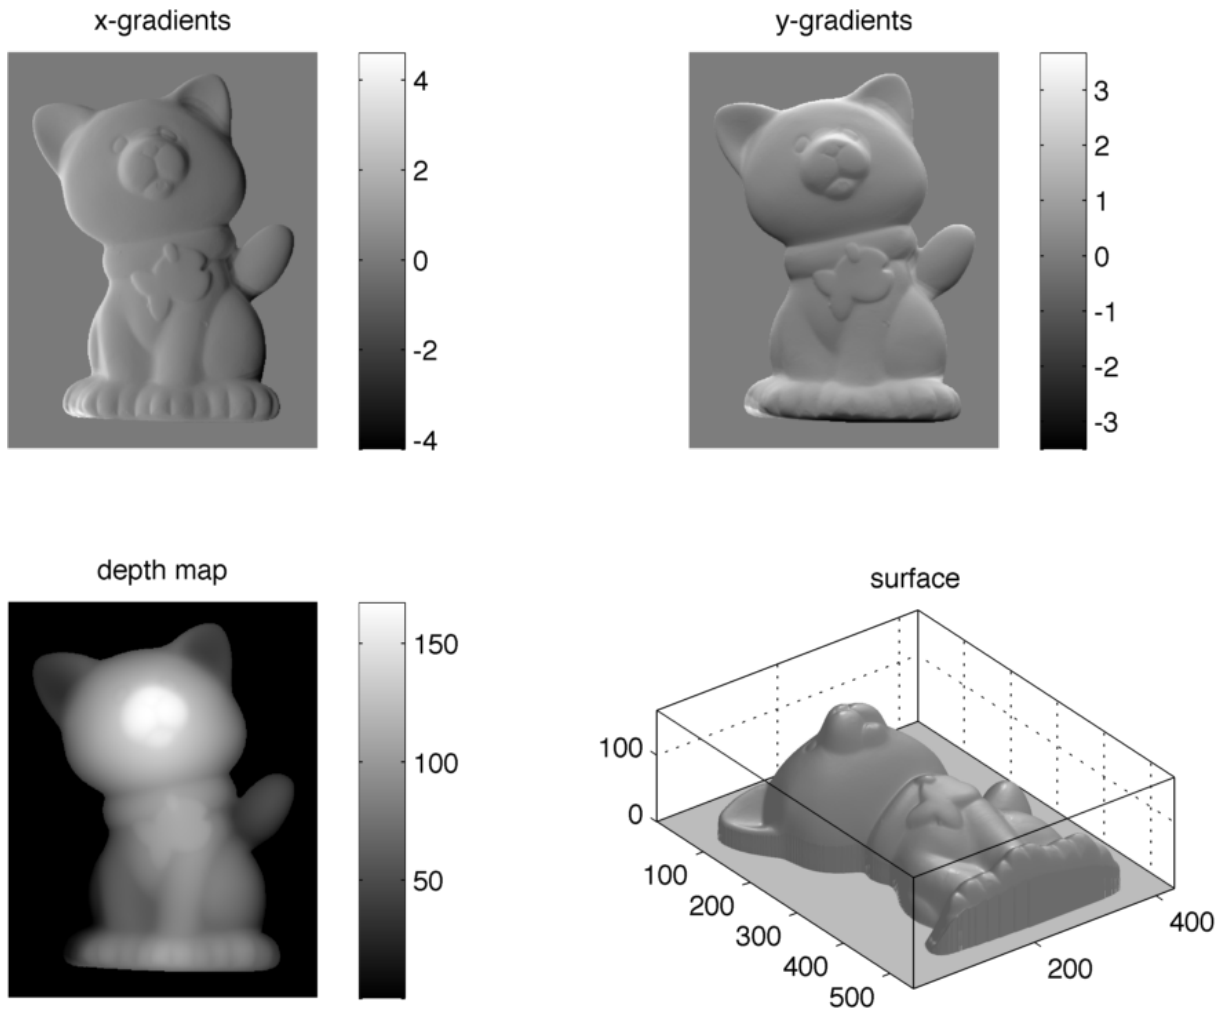


FIGURE 1. Photometric stereo reconstruction

Pr. 9.

(a) A possible Julia implementation is

```
function lsgd(A, b, mu, x0, nIters)
#
# Syntax:      x = lsgd(A, b, mu, x0, nIters)
```

```

#
# Inputs:      A is a m x n matrix
#
#              b is a vector of length m
#
#              mu is the step size to use, and must satisfy
#              0 < mu < 2 / norm(A)^2 to guarantee convergence
#
#              x0 is the initial starting vector (of length n) to use
#
#              nIters is the number of iterations to perform
#
# Outputs:     x is a vector of length n containing the approximate solution
#
# Description: Performs gradient descent to solve the least squares
#
#              \min_x \|b - A x\|_2
#
#
# Parse inputs
b = vec(b)
x0 = vec(x0)

# Gradient descent
x = x0
for _ in 1:nIters
    x -= mu * (A' * (A * x - b))
end

return x
end

```

- (b) Figure 2 shows a plot of $\|x_k - \hat{x}\|$ versus iteration k for one realization of the system with step size $\mu = 1/\sigma_1^2(A)$, for four values of noise standard deviation σ . Clearly the x_k iterates are converging to $\hat{x} = A^+b$.




FIGURE 2. Convergence of gradient descent for least squares problems with different noise levels.

Pr. 10.

- (a) A possible Julia implementation is

```

function lsngd(A, b, mu, x0, nIters)
#
# Syntax:      x = lsngd(A, b, mu, x0, nIters)
#
# Inputs:      A is a m x n matrix
#
#              b is a vector of length m
#
#              mu is the step size to use, and must satisfy
#              0 < mu < 1 / norm(A)^2 to guarantee convergence
#
#              x0 is the initial starting vector (of length n) to use
#
#              nIters is the number of iterations to perform
#
# Outputs:     x is a vector of length n containing the approximate solution
#
# Description: Performs Nesterov-accelerated gradient descent to solve the least
#              squares problem
#
#              \min_x \|b - Ax\|_2

    # Parse inputs
    b = vec(b)
    x0 = vec(x0)

    # Nesterov-accelerated gradient descent
    t = 0
    xLast = x0
    x = x0
    for _ = 1:nIters
        # t update
        tLast = t
        t = 0.5 * (1 + sqrt(1 + 4 * t^2))

        # z update
        z = x + ((tLast - 1) / t) * (x - xLast)

        # x update
        xLast = x
        x = z - mu * (A' * (A * z - b))
    end

    return x
end

```

- (b) Figure 3 compares sequences of $\|\mathbf{x}_k - \hat{\mathbf{x}}\|$ versus k for standard gradient descent and Nesterov-accelerated gradient descent.
- (c) For step sizes $\mu = \{0.25, 0.5, 0.75, 1\} / \sigma_1^2(A)$, clearly Nesterov-accelerated gradient descent converges faster than standard gradient descent, although the convergence is not necessarily monotone.




FIGURE 3. Nesterov-accelerated vs. standard gradient descent for a least squares problem, for four values of step size μ .